

SUMESH R -20104169

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv("1_fiat500_VehicleSelection_Dataset.csv")[0:1500].dropna(axis="columns",
df
```

```
Out[2]:
```

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | p |
|------|--------|--------|--------------|-------------|----------|-----------------|-----------|-------------|-----|
| 0 | 1.0 | lounge | 51.0 | 882.0 | 25000.0 | 1.0 | 44.907242 | 8.611559868 | € |
| 1 | 2.0 | pop | 51.0 | 1186.0 | 32500.0 | 1.0 | 45.666359 | 12.24188995 | € |
| 2 | 3.0 | sport | 74.0 | 4658.0 | 142228.0 | 1.0 | 45.503300 | 11.41784 | 4 |
| 3 | 4.0 | lounge | 51.0 | 2739.0 | 160000.0 | 1.0 | 40.633171 | 17.63460922 | € |
| 4 | 5.0 | pop | 73.0 | 3074.0 | 106880.0 | 1.0 | 41.903221 | 12.49565029 | 5 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 1495 | 1496.0 | pop | 62.0 | 3347.0 | 80000.0 | 3.0 | 44.283878 | 11.88813972 | 7 |
| 1496 | 1497.0 | pop | 51.0 | 1461.0 | 91055.0 | 3.0 | 44.508839 | 11.46907997 | 7 |
| 1497 | 1498.0 | lounge | 51.0 | 397.0 | 15840.0 | 3.0 | 38.122070 | 13.36112022 | 10 |
| 1498 | 1499.0 | sport | 51.0 | 1400.0 | 60000.0 | 1.0 | 45.802021 | 9.187789917 | 10 |
| 1499 | 1500.0 | pop | 51.0 | 1066.0 | 53100.0 | 1.0 | 38.122070 | 13.36112022 | € |

1500 rows × 9 columns



```
In [3]: df.head()
```

```
Out[3]:
```

| | ID | model | engine_power | age_in_days | km | previous_owners | lat | lon | price |
|---|-----|--------|--------------|-------------|----------|-----------------|-----------|-------------|-------|
| 0 | 1.0 | lounge | 51.0 | 882.0 | 25000.0 | 1.0 | 44.907242 | 8.611559868 | 8900 |
| 1 | 2.0 | pop | 51.0 | 1186.0 | 32500.0 | 1.0 | 45.666359 | 12.24188995 | 8800 |
| 2 | 3.0 | sport | 74.0 | 4658.0 | 142228.0 | 1.0 | 45.503300 | 11.41784 | 4200 |
| 3 | 4.0 | lounge | 51.0 | 2739.0 | 160000.0 | 1.0 | 40.633171 | 17.63460922 | 6000 |
| 4 | 5.0 | pop | 73.0 | 3074.0 | 106880.0 | 1.0 | 41.903221 | 12.49565029 | 5700 |

Data cleaning and pre processing

In [4]:

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 9 columns):
#   Column              Non-Null Count  Dtype
---  -
0   ID                   1500 non-null   float64
1   model                1500 non-null   object
2   engine_power         1500 non-null   float64
3   age_in_days          1500 non-null   float64
4   km                   1500 non-null   float64
5   previous_owners      1500 non-null   float64
6   lat                  1500 non-null   float64
7   lon                  1500 non-null   object
8   price                1500 non-null   object
dtypes: float64(6), object(3)
memory usage: 105.6+ KB
```

In [5]:

`df.describe()`

Out[5]:

| | ID | engine_power | age_in_days | km | previous_owners | lat |
|--------------|-------------|--------------|-------------|---------------|-----------------|-------------|
| count | 1500.000000 | 1500.000000 | 1500.000000 | 1500.000000 | 1500.000000 | 1500.000000 |
| mean | 750.500000 | 51.875333 | 1641.629333 | 53074.900000 | 1.126667 | 43.545904 |
| std | 433.157015 | 3.911606 | 1288.091104 | 39955.013731 | 0.421197 | 2.112907 |
| min | 1.000000 | 51.000000 | 366.000000 | 1232.000000 | 1.000000 | 36.855839 |
| 25% | 375.750000 | 51.000000 | 670.000000 | 20000.000000 | 1.000000 | 41.802990 |
| 50% | 750.500000 | 51.000000 | 1035.000000 | 38720.000000 | 1.000000 | 44.360376 |
| 75% | 1125.250000 | 51.000000 | 2616.000000 | 78170.250000 | 1.000000 | 45.467960 |
| max | 1500.000000 | 77.000000 | 4658.000000 | 235000.000000 | 4.000000 | 46.795612 |

In [6]:

`df.columns`

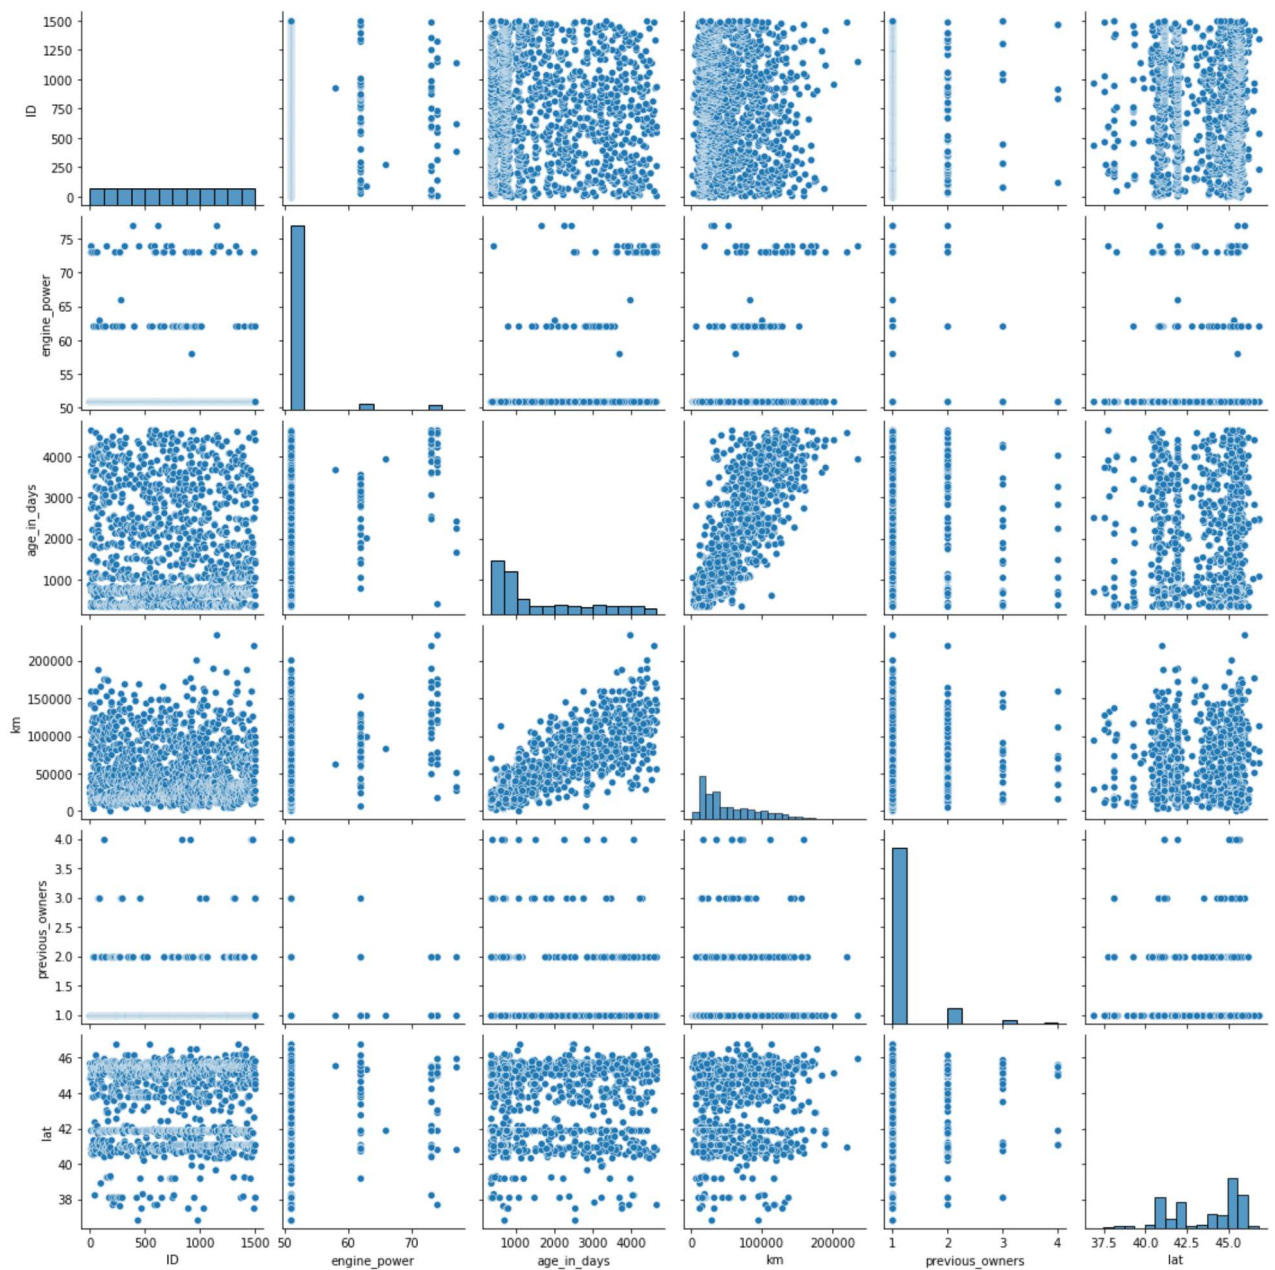
```
Out[6]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
              'lat', 'lon', 'price'],
              dtype='object')
```

EDA and VISUALIZATION

In [7]:

`sns.pairplot(df)`


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x26b9813ad90>
```



```
In [8]: sns.distplot(df["engine_power"])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)


```
Out[8]: <AxesSubplot:xlabel='engine_power', ylabel='Density'>
```



```
In [9]: df1 = df[['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
                'lat', 'lon', 'price']]
```

```
In [10]: sns.heatmap(df1.corr())
```

Out[10]: <AxesSubplot:>



```
In [11]: x = df1[['ID', 'age_in_days', 'km', 'previous_owners',
                'lat']]
          y = df1['engine_power']
```

split the data into training and test data

```
In [12]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: lr = LinearRegression()  
         lr.fit(x_train, y_train)
```

```
Out[13]: LinearRegression()
```

```
In [14]: lr.intercept_
```

```
Out[14]: 52.76671328218245
```


```
In [15]: coeff = pd.DataFrame(lr.coef_, x.columns, columns = ['Co-efficient'])  
         coeff
```

```
Out[15]:
```

| | Co-efficient |
|------------------------|--------------|
| ID | -0.000219 |
| age_in_days | 0.000745 |
| km | 0.000005 |
| previous_owners | -0.581619 |
| lat | -0.036866 |

```
In [16]: prediction = lr.predict(x_test)  
         plt.scatter(y_test, prediction)
```

```
Out[16]: <matplotlib.collections.PathCollection at 0x26b9a027e20>
```



```
In [17]: lr.score(x_test, y_test)
```

```
Out[17]: 0.107544924049784
```